

C* Primer

Part 01: Cassandra and Modeling Basics

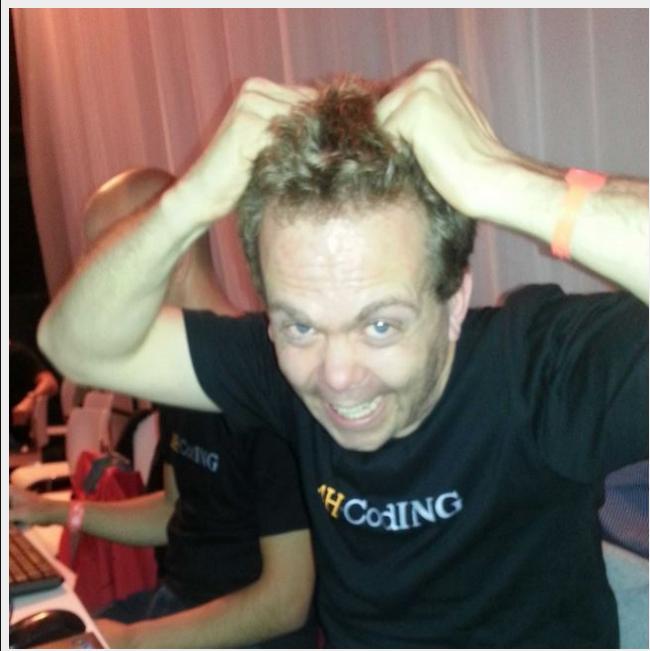
```
SELECT * FROM presenters WHERE name IN ('Christopher Reedijk', 'Gary Stewart');  
name          | title        | company | area | twitter  
-----+-----+-----+-----+-----  
Christopher Reedijk | Dev Engineer | ING      | NL    | @creedijk  
Gary Stewart       | Dev Engineer | ING      | NL    | @Gaz_GandA
```



ING Nederland ([@ingnl](#))

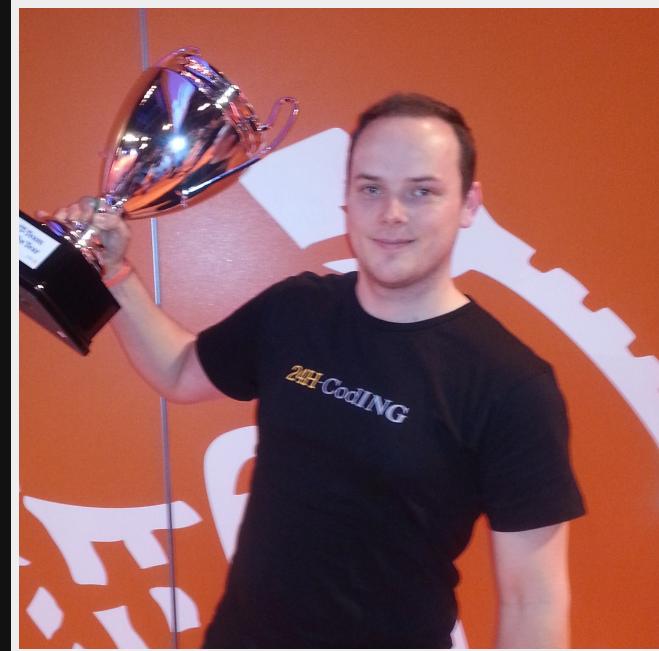
This presentation is based on a 101 presentation by
Hayato Shimizu ([@hayato_shimizu](#))

Short introduction



Gary Stewart

- Dev Engineer
- Love/hate relationship with C*
- @ING since: 01-01-2012
- aka Legacy



Christopher Reedijk

- Dev Engineer
- C* enthusiast
- @ING since: 01-01-2011
- aka Grumpy

3 Challenges

challenge 1

Improve **availability**
without trading **consistency**

availability & consistency

Consequences of not being available



Problemen bij internetbankieren ING

AMSTERDAM - Internetbankieren kampert donderdagmiddag weer met een kleine storing. Een fractie van de klanten heeft problemen met inloggen of het...



Internetbankieren ING beperkt toegankelijk

AMSTERDAM - Klanten van ING kunnen woensdag soms geen gebruik maken van internetbankieren omdat de bank een maximum aan verkeer heeft ingesteld.



Internetbankieren ING onbereikbaar door onderhoud

AMSTERDAM - Vanwege uitgelopen onderhoud aan internetbankieren bij ING was de dienst dinsdagochtend niet bereikbaar. Inmiddels kunnen klanten...

source: [nu.nl](#)

availability & consistency

Consequences of not being consistent



challenge 2

Aim to be **easier** scalable
Changes are happening at an
increasing pace

easier scalable

Stop focusing on the **expected load**



yesterday

source: bradfrostweb.com

easier scalable

Start focusing on the unexpected load

THIS IS THE WEB.



today

THIS WILL BE THE WEB.



tomorrow

source: bradfrostweb.com

challenge 3

Adopt **new ways of thinking**
Become the top **engineering**
company

new way of thinking

ING's **culture** is changing fast.
Waterfall to Scrum to DevOps
in less than **2 years**

new way of thinking

Aim to approach problems as **green field** to understand the essence
"Un-learn" principles that don't scale

Use **cache** correctly!

Reduce **locking** (transaction)

Apache Cassandra

Created by Avinash Lakshman and Prashant Malik at Facebook

Cassandra is

- A distributed database
- Highly scalable
- Fault tolerant
- Consistency is tunable
- High throughput
- Network topology aware
- Multi-DC Active-Active
- Written in Java
- A column family
- CQL3

Cassandra is not

- ACID compliant database
- A relational database
- ANSI SQL compliant
- A document database

DataStax and Cassandra



Founded in 2010 by Jonathan Ellis and Matt Pfeil

80% Apache Cassandra code contribution

Offer commercial support for DataStax Enterprise version of Cassandra

DataStax Enterprise integrates Search and Analytics

Head Quarter in San Francisco Bay area

EMEA office opened in March 2013

Native Protocol / DataStax Drivers

CQL3 Only

Fully asynchronous protocol, using Netty
Server notifications

Java, C#, Python, node.js and Ruby so far, more will be added
Many policies, including TokenAwarePolicy,
DowngradeConsistencyRetryPolicy, etc...
Metrics (metrics.codahale.com) included

Maven:

```
<dependency>
    <groupid>com.datastax.cassandra</groupid>
    <artifactid>cassandra-driver-core</artifactid>
    <version>2.0.2</version>
</dependency>
```

Cassandra Internals

Tunable Consistency

**Configureable Consistency Level (CL)
per read and write action**

Writes

- Any
- One
- Two
- Three
- Quorum
- Local_Quorum
- Each_Quorum
- All

Reads

- One
- Two
- Three
- Quorum
- Local_Quorum
- Each_Quorum
- All

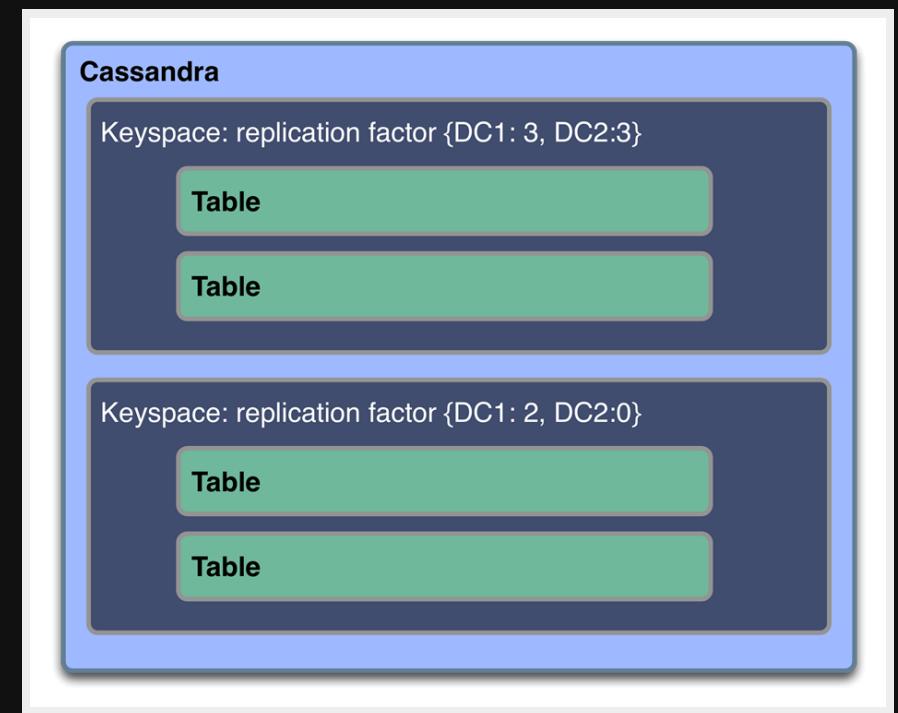
Logical Grouping

Keyspace: group of tables

Table: group of data

Replication Factor:

Number of copies of data



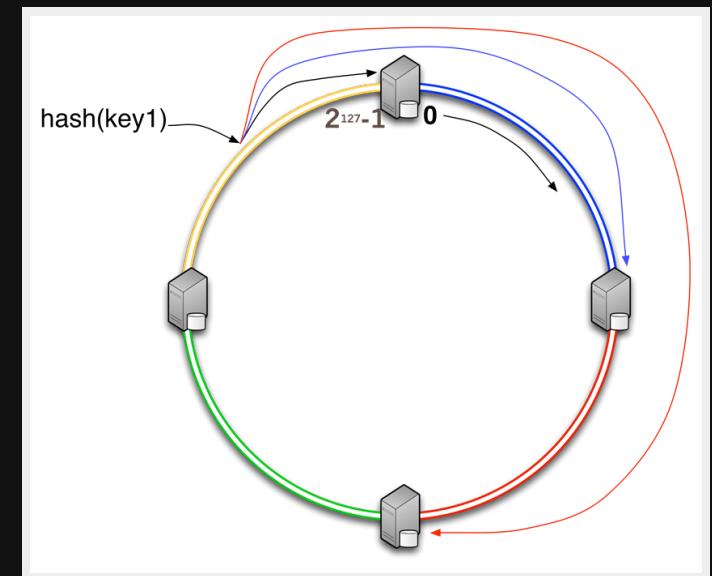
Cassandra Replication Strategy

Token Range 0 -> $2^{127}-1$ in Ring Formation

Consistent Hashing Algorithm

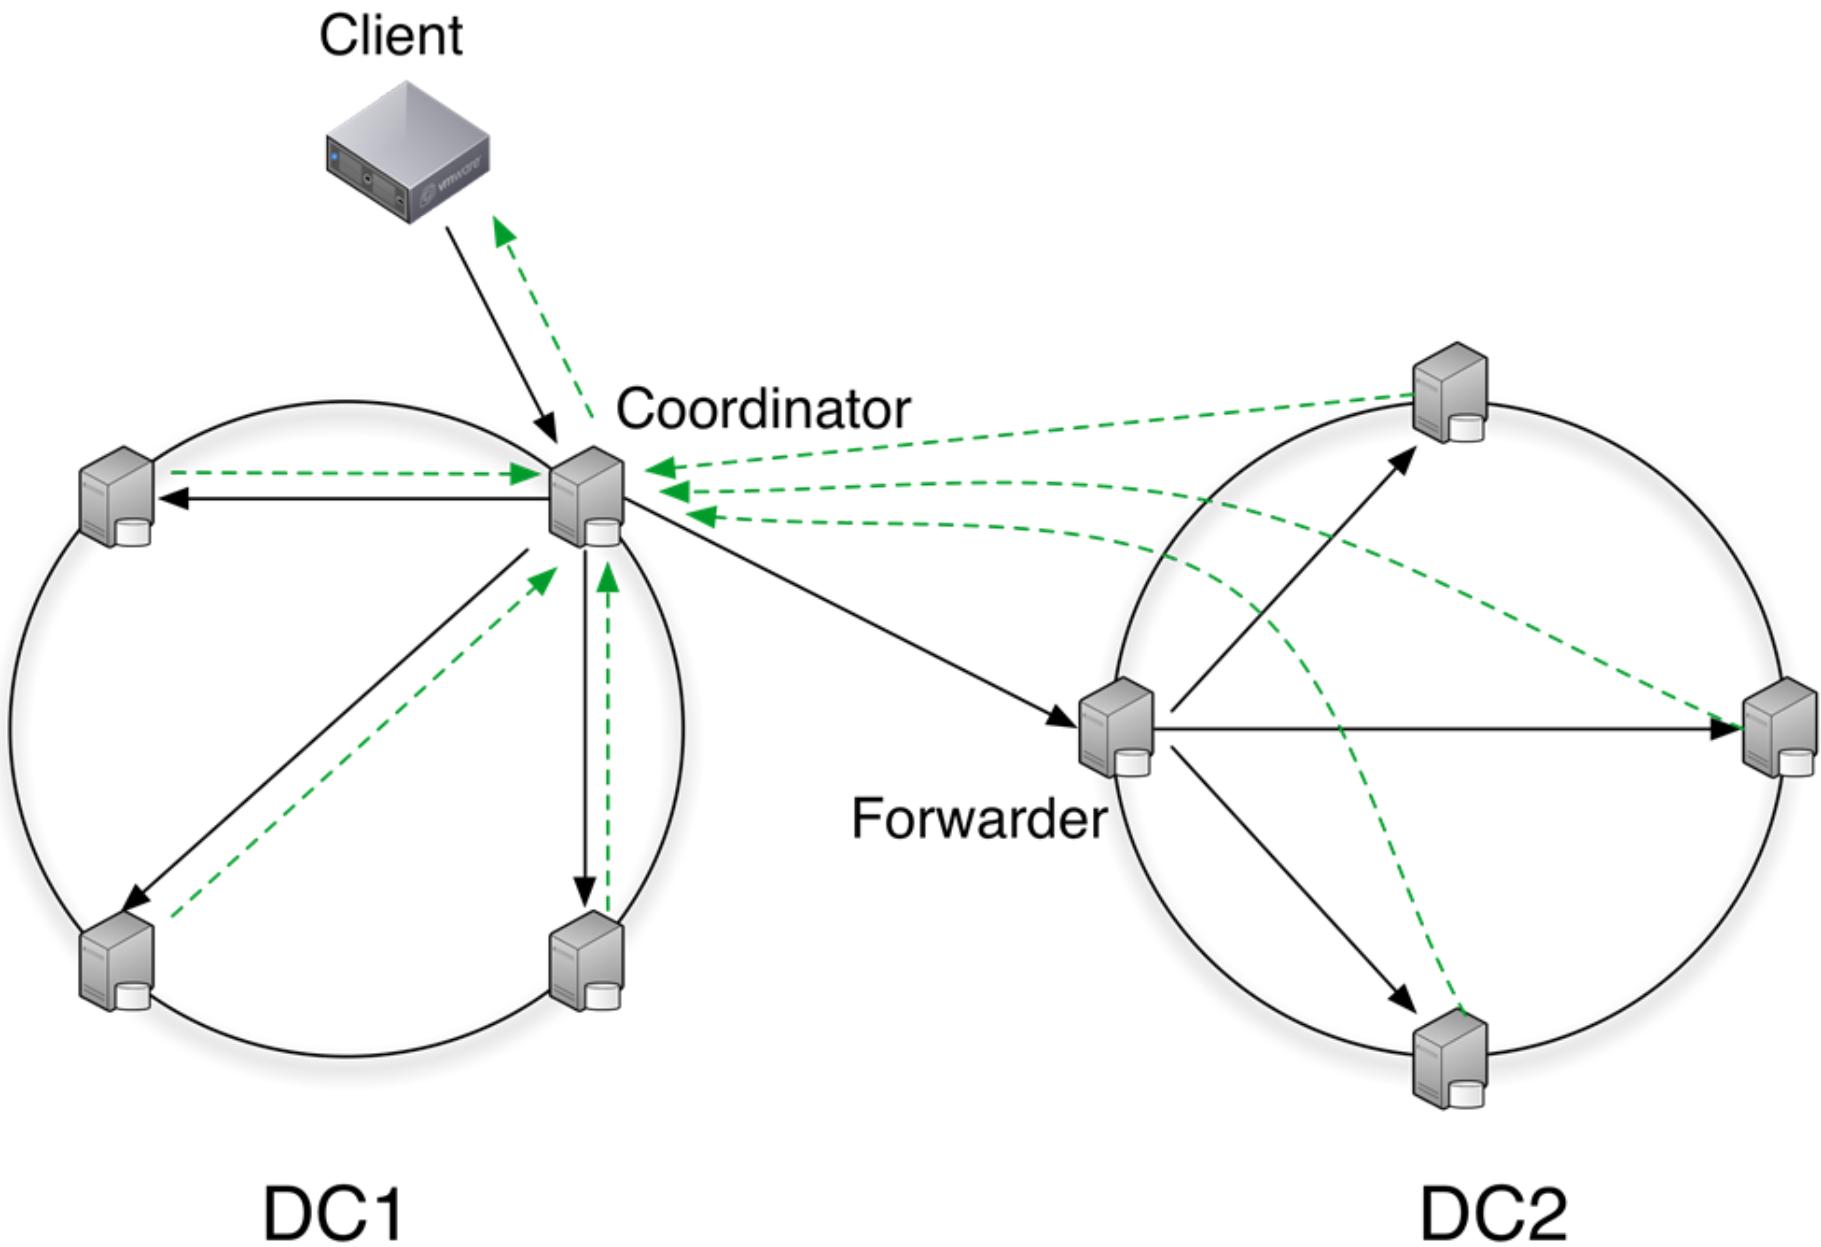
Replica nodes in clockwise

Gossip protocol between nodes



Replication Factor (RF) = 3

Network Topology Awareness

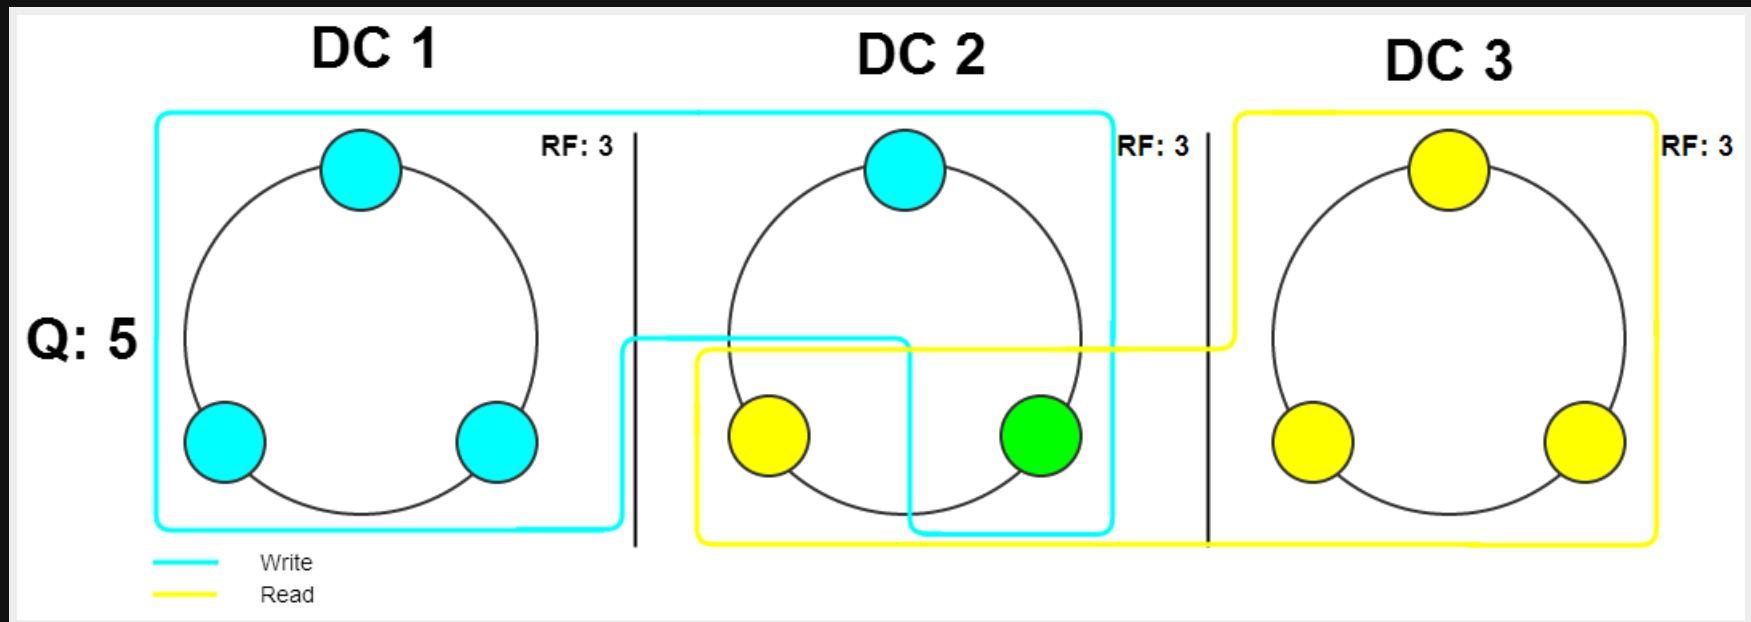


Replication Factor and Data Consistency In a Distributed System

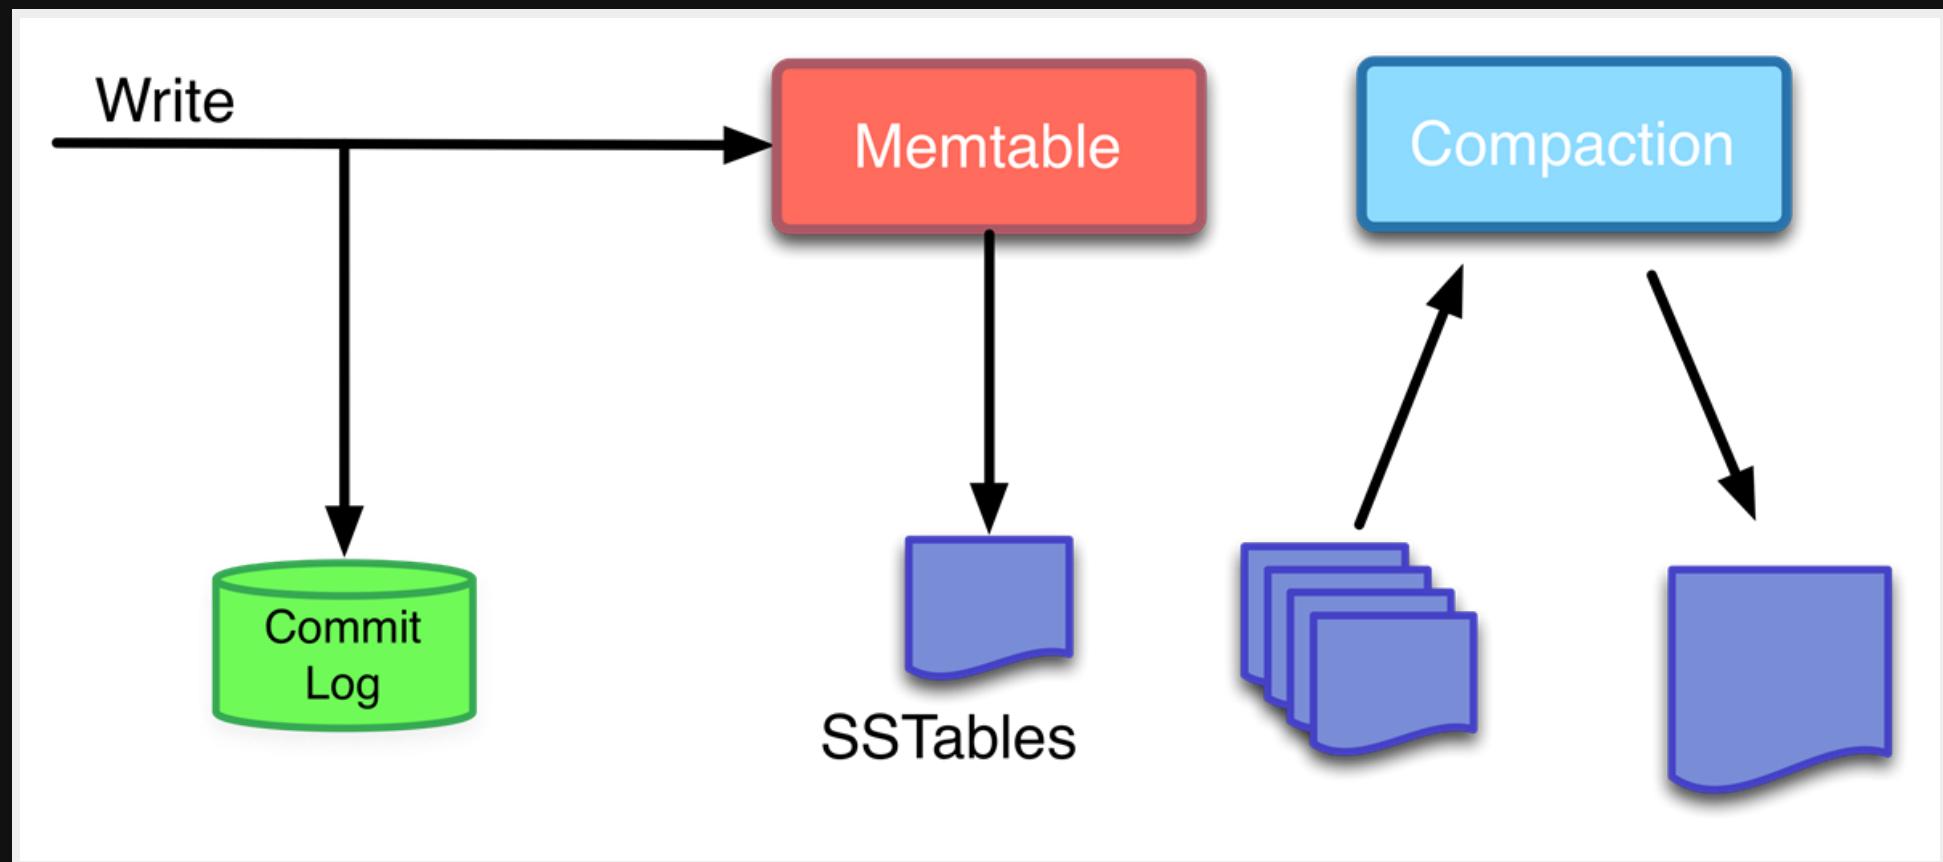
CAP – Consistency, Availability, Partition Tolerance

Consistency is achieved with $W + R > RF$

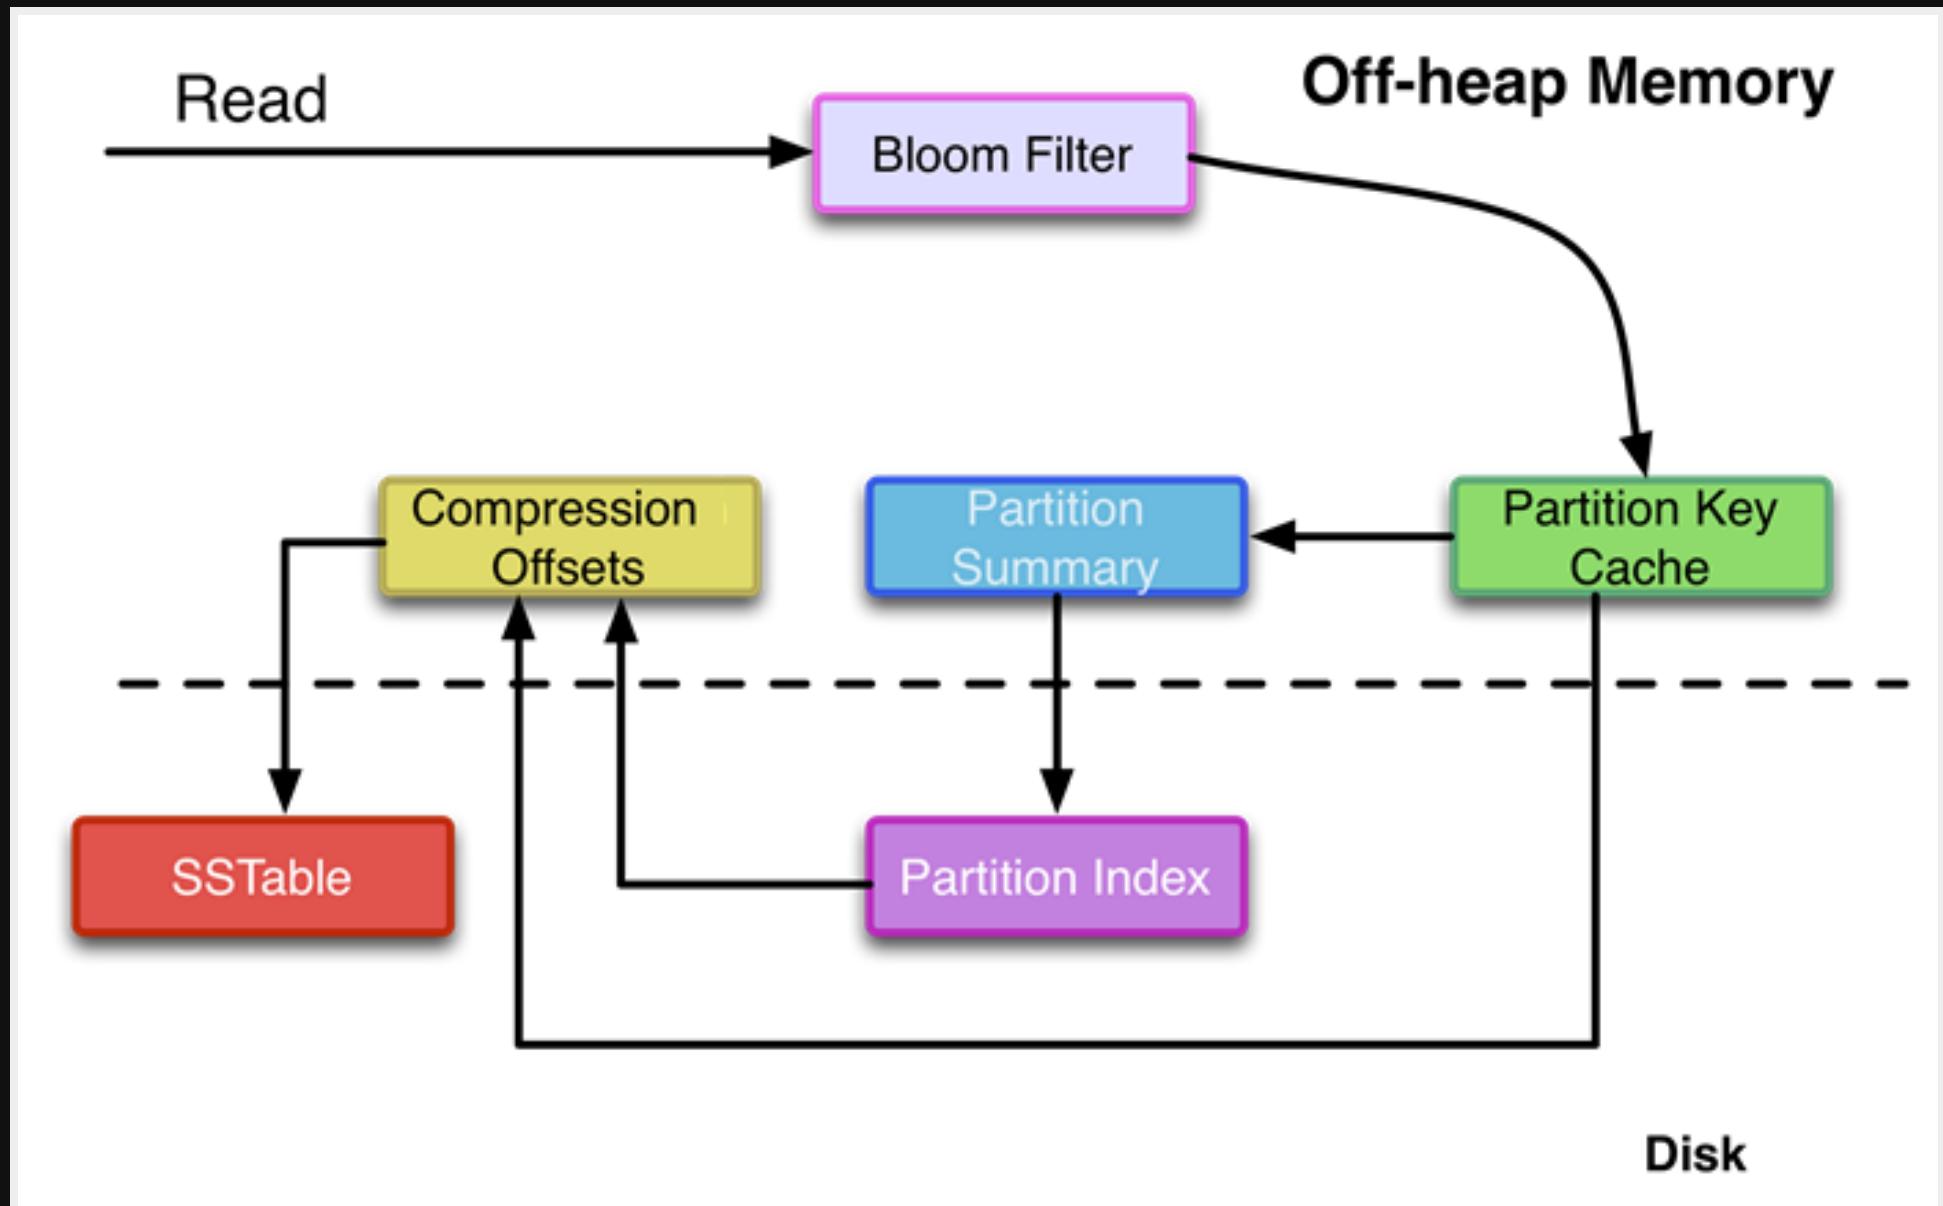
$$\text{Quorum} = (\text{Integer}) \frac{RF}{2} + 1$$



Write Path



Read Path



CQL3

Why CQL3?

Excellent language for modeling data structures

Usability

Readability

Familiarity

Programming language agnostic

Apache Thrift, Hector, Astyanax days

```
Column col =
    new Column(ByteBuffer.wrap("product".getBytes()))
);
col.setValue(ByteBuffer.wrap("phone".getBytes()));
col.setTimestamp(System.currentTimeMillis());

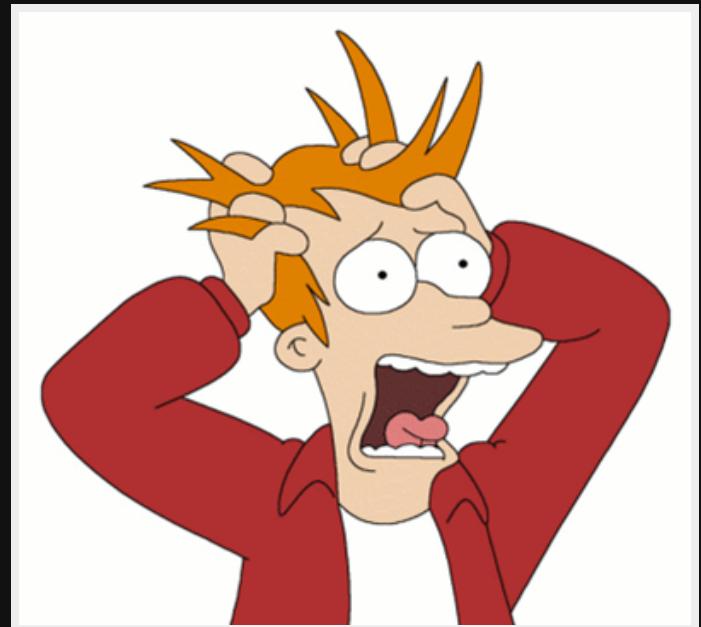
ColumnOrSuperColumn c = new ColumnOrSuperColumn();
c.setColumn(col);

Mutation m = new Mutation();
m.setColumn_or_supercolumn(c);

List ms = new ArrayList();
ms.add(m);

Map> cf =
    new HashMap();
cf.put("products", m);

Map>> records =
    new HashMap>>();
records.put(ByteBuffer.wrap("key".getBytes()), cf);
client.batch_mutate(records, consistencyLevel);
```



CQL3

```
INSERT INTO products (product_id, product)
VALUES ("id12321", "phone");
```



Data Modelling

Primary Key

The Primary Key

- The key that uniquely identifies a row
- A primary key consists of:
 - A (composite) **partition key**
 - One or more **clustering columns**
 - e.g. **PRIMARY KEY (partition key, cluster columns, ...)**
- The **partition key** determines on which node the partition resides
- Data is ordered in **cluster column** order within the partition

Creating a table

```
CREATE TABLE cities (
    city_name    varchar,
    elevation    int,
    population   int,
    latitude     float,
    longitude    float,
    PRIMARY KEY (city_name)
);
```

We can visualize it this way:

city_name	elevation	population	latitude	longitude
Springfield	978	60608	39°55'37"	83°48'15" W

In this example:
partition key = primary key

Composite partition key

```
CREATE TABLE cities (
    city_name  varchar,
    state varchar,
    elevation float,
    latitude float,
    longitude float
    population float,
    PRIMARY KEY
        ( (city_name,state) )
);
```

(city_name, state)	elevation	latitude	longitude	population
Springfield, OH	978	39.926	-83.804	60608
Springfield, IL	39.783	39.783	-89.650	116250

Each city gets its own partition!

Clustering columns

```
CREATE TABLE sporty_league (
    team_name  varchar,
    player_name varchar,
    jersey     int,
    PRIMARY KEY (team_name,
                  player_name)
);
```

team_name Springers	Adler	Bélanger	Foote
team_name Mighty Mutts	Buddy	Lucky	
team_name Peppers	Aaron	Baker	Cabrera

Not ordered

Ordered by player_name

Simple Select

```
SELECT *  
FROM sporty_league;
```

team_name	player_name	jersey
Peppers	Aaron	17
Peppers	Baker	62
Peppers	Cabrera	25
Springers	Adler	86
Springers	Bélanger	13
Springers	Foote	99
Mighty Mutts	Buddy	32
Mighty Mutts	Lucky	7

Not ordered

The diagram illustrates a transformation from an initial state to a final state. On the left, a vertical double-headed arrow is labeled "Not ordered". To its right is a table with three rows. The first row has orange header cells for "team_name", "Adler", "Bélanger", and "Foote". The second row has teal header cells for "team_name", "Buddy", and "Lucky". The third row has orange header cells for "team_name" and "Peppers". A red arrow points from the right side of the initial table to the right side of the final table, which is labeled "Ordered by player_name". The final table has four columns: "team_name", "Adler" (orange), "Bélanger" (teal), and "Foote" (teal). The data rows are: Row 1 (orange) has "Springers" in the first column and "86" in the second column; Row 2 (teal) has "Mighty Mutts" in the first column and "32" in the second column; Row 3 (orange) has "Peppers" in the first column and "17" in the second column.

team_name	Adler	Bélanger	Foote
Springers	86	13	99
team_name	Buddy	Lucky	
Mighty Mutts	32	7	
team_name	Aaron	Baker	Cabrera
Peppers	17	62	25

"Table scan" - slow! (Limited to 10,000 rows by default)
Use LIMIT keyword to choose fewer or more rows

Simple Select on Partition Key and Cluster Columns

```
SELECT *  
  FROM sporty_league  
 WHERE team_name = 'Mighty Mutts';
```

team_name	player_name	jersey
Mighty Mutts	Buddy	32
Mighty Mutts	Lucky	7

```
SELECT *  
  FROM sporty_league  
 WHERE team_name = 'Mighty Mutts'  
   AND player_name = 'Lucky';
```

team_name	player_name	jersey
Mighty Mutts	Lucky	7

Insert/Update

```
INSERT INTO sporty_league (team_name, player_name, jersey)
VALUES ('Mighty Mutts', 'Felix', 90);
```

```
UPDATE sporty_league
SET jersey = 100
WHERE team_name = 'Mighty Mutts'
AND player_name = 'Felix';
```

Data Types

CQL Type	Constants	Description
ascii	strings	US-ASCII character string
bigint	integers	64-bit signed long
blob	blobs	Arbitrary bytes (no validation), expressed as hexadecimal
boolean	booleans	true or false
counter	integers	Distributed counter value (64-bit long)
decimal	integers, floats	Variable-precision decimal
double	integers	64-bit IEEE-754 floating point
float	integers, floats	32-bit IEEE-754 floating point
inet	strings	IP address string in IPv4 or IPv6 format ^[1]
int	integers	32-bit signed integer
list	n/a	A collection of one or more ordered elements
map	n/a	A collection of one or more timestamp, value pairs
set	n/a	A collection of one or more elements
text	strings	UTF-8 encoded string
timestamp	integers, strings	Date plus time, encoded as 8 bytes since epoch
uuid	uuids	Type 1 or type 4 UUID in standard UUID format
timeuuid	uuids	Type 1 UUID only (CQL 3)
varchar	strings	UTF-8 encoded string
varint	integers	Arbitrary-precision integer

Collections

CQL supports having columns that contain collections of data

The collection types include:
Set, List and Map
(64K max limitation)

```
CREATE TABLE collections_example (
    id int PRIMARY KEY,
    set_example set<text>,
    list_example list<text>,
    map_example map<int, text>
);
```

These data types are intended to support the type of 1-to-many relationships that can be modeled in a relational DB e.g. a user has many email addresses

Some performance considerations around collections

- Often more efficient to denormalise further rather than use collections if intending to store lots of data.
- Favour sets over list – lists not very performant

Performance considerations

The best queries are in a single partition.

i.e. WHERE partition key =

Queries that span multiple partitions are **s-l-o-w**

i.e. WHERE partition key IN (, , ...);

Queries that span multiple cluster columns are **fast**

i.e. WHERE partition key = and clustering key =

(Clustering) Order By

```
SELECT *
  FROM sporty_league
 WHERE team_name = 'Mighty Mutts'
 ORDER BY player_name DESC;
```

```
CREATE TABLE sporty_league (
    team_name    varchar,
    player_name  varchar,
    jersey      int,
    PRIMARY KEY  (team_name, player_name)
) WITH CLUSTERING ORDER BY
      (player_name DESC);
```

Partition keys are not ordered!

You can only order by a clustering column

Data will be ordered ASCending by default

You can also use the ORDER BY keyword, but only on the clustering column(s), which don't have CLUSTERING ORDER BY defined!

De-normalize

Don't do this

```
CREATE TABLE users (
    id uuid PRIMARY KEY,
    username text,
    email text,
    street_address text,
    town text,
);

CREATE INDEX user_street ON
    users (street_address);
CREATE INDEX user_town ON
    users (town);

SELECT *
    FROM users
    WHERE town = 'Stockholm'
        AND street_address = 'Sankt Eriksgatan'
    ALLOW FILTERING;
```

Do this instead

```
CREATE TABLE users (
    id uuid PRIMARY KEY,
    username text,
    email text,
    street_address text,
    town text,
);

CREATE TABLE user_street (
    town text,
    street text,
    user_id uuid,
    PRIMARY KEY (
        (town, street), user_id)
);
```

Time to Live (TTL)

```
INSERT INTO users (id, first, last)
VALUES ('abc123', 'abe', 'lincoln')
USING TTL 3600; // Expires data in one hour
```

Time Series

```
CREATE TABLE sensor_data (
    sensor_id int,
    date varchar,
    time timeuuid,
    value int,
    PRIMARY KEY( (sensor_id, date), time ) )
WITH CLUSTERING ORDER BY (time DESC);
```

```
INSERT INTO sensor_data
    (sensor_id, date, time, value)
VALUES (1, '20130908',
        1883-11e3-8ffd-0800200c9a66, 432)
USING TTL 604800;
```

```
SELECT time, value, TTL (value)
  FROM sensor_data
 WHERE sensor_id = 1 AND date = '20130908'
   AND time > maxTimeuuid('2013-09-08 10:05+0000')
   AND time < minTimeuuid('2013-09-08 15:00+0000');
```

time		value		ttl(value)
1883-11e3-8ffd-0800200c9a66		432		604257
1887-11e3-8ffd-0800200c9a66		433		604452

Lightweight Transactions

Introduced in Cassandra 2.0

Example:

```
INSERT INTO customer_account (customerID, customer_email)
VALUES ('LauraS', 'lauras@gmail.com')
    IF NOT EXISTS;

UPDATE customer_account
    SET customer_email='laurass@gmail.com'
    IF customer_email='lauras@gmail.com';
```

Great for 1% of your application – but not recommended to be used too much!

Tracing

You can turn on tracing on or off for queries with the TRACING ON | OFF command

This can help you understand what Cassandra is doing and identify any performance problems

```
cqlsh:ecom> SELECT vendor, order_id, user_id, quantity, total_cost, product_id, product_name, order_timestamp FROM order_by_vendor WHERE vendor='YooDoo BBQ & Grill Franchising' AND bucket = 1;
vendor          | order_id | user_id | quantity | total_cost | product_id | product_name          | order_timestamp
YooDoo BBQ & Grill Franchising | 0235 | U1949 | 8 | 119.68 | P1632 | Sobe - Cranberry Grapefruit | 2013-08-08 23:02:50+0000

Tracing session: 9d316d90-4743-11e3-bda5-1166498cf1d9
activity
execute_cql3_query | timestamp | source | source_elapsed
-----+-----+-----+-----+
Parsing SELECT vendor, order_id, user_id, quantity, total_cost, product_id, product_name, order_timestamp FROM order_by_vendor WHERE vendor='YooDoo BBQ & Grill Franchising' AND bucket = 1 LIMIT 10000; | 08:29:07,691 | 192.168.184.176 | 0
Preparing statement | 08:29:07,691 | 192.168.184.176 | 140
Executing single-partition query on order_by_vendor | 08:29:07,692 | 192.168.184.176 | 350
Acquiring sstable references | 08:29:07,692 | 192.168.184.176 | 1005
Merging memtable tombstones | 08:29:07,692 | 192.168.184.176 | 1120
Merging data from memtables and 0 sstables | 08:29:07,692 | 192.168.184.176 | 1155
Read 1 live and 0 tombstoned cells | 08:29:07,692 | 192.168.184.176 | 1354
Request complete | 08:29:07,692 | 192.168.184.176 | 1680
```

<http://www.datastax.com/dev/blog/tracing-in-cassandra-1-2>

all challenges appear again...

Improve availability

Aim to be scalable

Adopt new ways of thinking

3 Facts

fact 1

Availability is easier

availability is easier, because

Masterless architecture

No outages in contrast to master-slave architecture

This does make C* chatty

availability is easier, because

Replication of data

Replication factor is **configurable** per keyspace

Increasing nodes **does not** mean higher availability

Increasing the replication factor **does**

fact 2

Performance is easier

performance is easier, and

Solve your **reads** with your **writes**

De-normalization is acceptable for solving reads

performance is easier, but

Know your **partitions**

Large partitions will hurt!

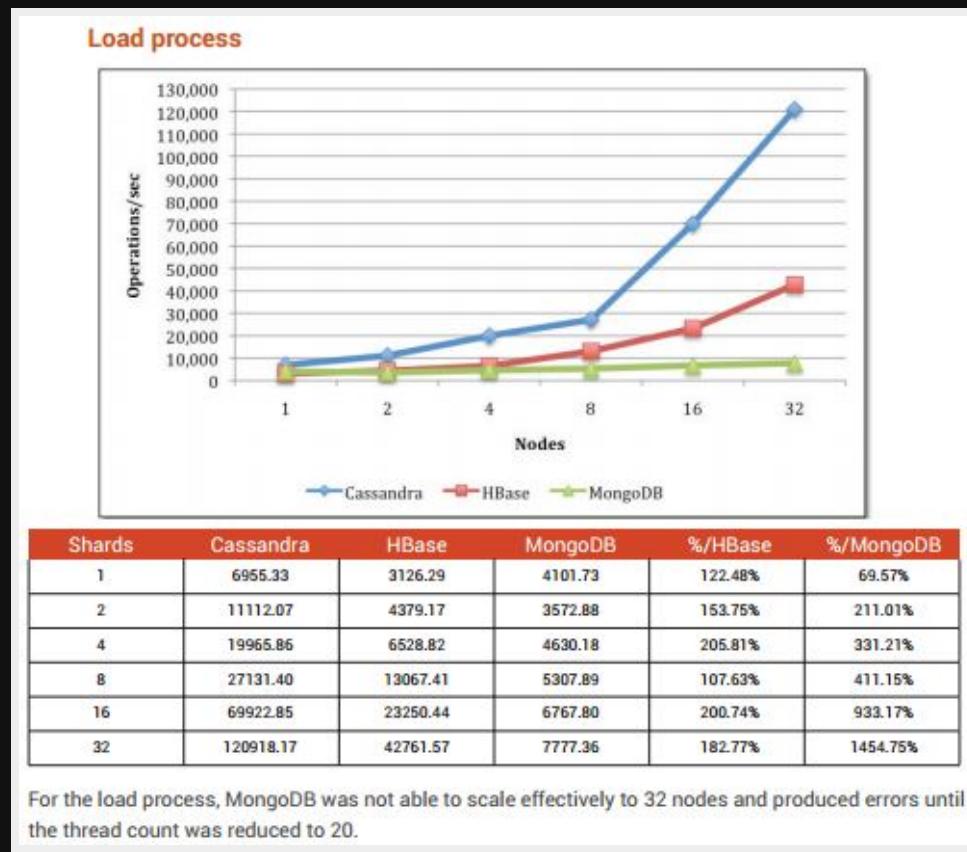
Partition Key	Clustering Key	Col A	Col B
partitionKey_1	1	foo	...
	2	bar	...
partitionKey_2	1	foo	...
	2	bar	...
	3	baz	...

	1000000	qux	...

performance is easier, because

Linear scalable

Adding nodes will increase your throughput through **sharding**



source: [Datastax](#)

fact 3

Consistency is harder

consistency is harder, and

Application developer has much more responsibility

Focus on the **data model** and the **flow** of data is key

noSQL versus **SQL**

consistency is harder, because

CAP theorem

Consistency (all nodes see the same data at the same time)

Availability (a guarantee that every request receives a response about whether it was successful or failed)

Partition tolerance (the system continues to operate despite arbitrary message loss or failure of part of the system)

consistency is harder, because

Transactions and queuing are **anti-patterns**

C* v2.0 reduces the complexity by using lightweight transactions
(paxos)

all facts appear...

Availability is easier | challenge 1: availability & consistency

Performance is easier | challenge 2 easier scalable

Consistency is harder | challenge * availability & consistency, easier
scalable, new way of thinking

all challenges covered...

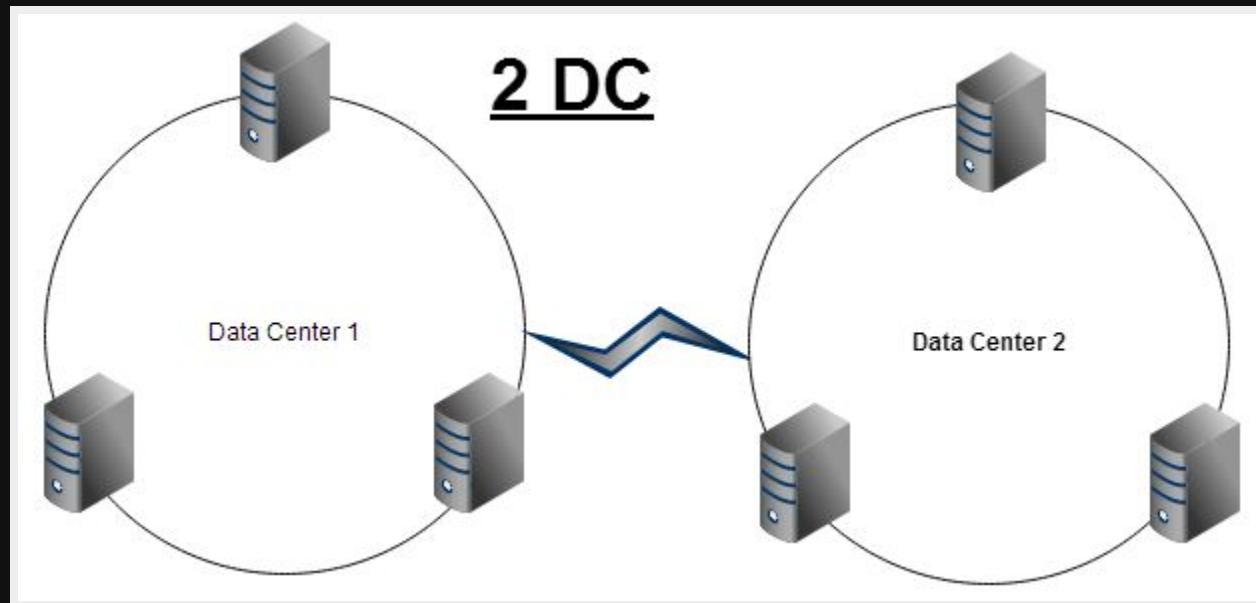
C*@ING

project	use case	status
Betelgeuze	Cache of Customer Data (KRO)	September 2014 in PRD
OI: Financial Fit	SOR for soft data	September 2014 in PRD
RTPE	New payments engine for on-us payments	February 2015 in PRD
Faster Than Light	Cache of Customer Data (MDM)	PoC
Ideal, Sales Support, NGINX, ...	Session Management	In DEV
Ideal	Order Management	Product Backlog
Credit Cards	Cache external data	Awaiting budget
Availability Dashboard	Time series	In DEV
Realtime Account Forecasting	Cache of SAM data	PoC

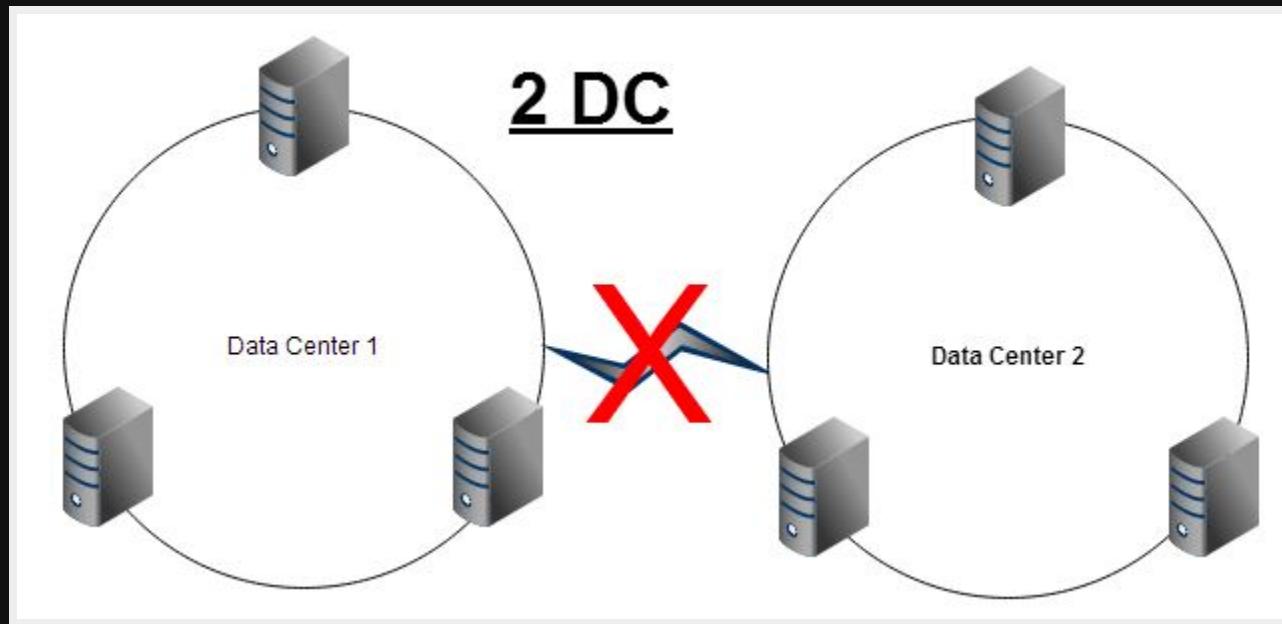
3 Questions - Bonus

question 1

What **risks** do we have with 2 DC's and using Write/Read CL of **LOCAL_QUORUM**?



Split-brain



Who wins? **Neither** side knows the other side!

CAP theorem - choose wisely

Use QUORUM but **risk availability**
when there are issues with the WAN link

Arrange a third DC... **€€€**

Applications apply DC stickiness
whilst processing data which **increases complexity**

question 2

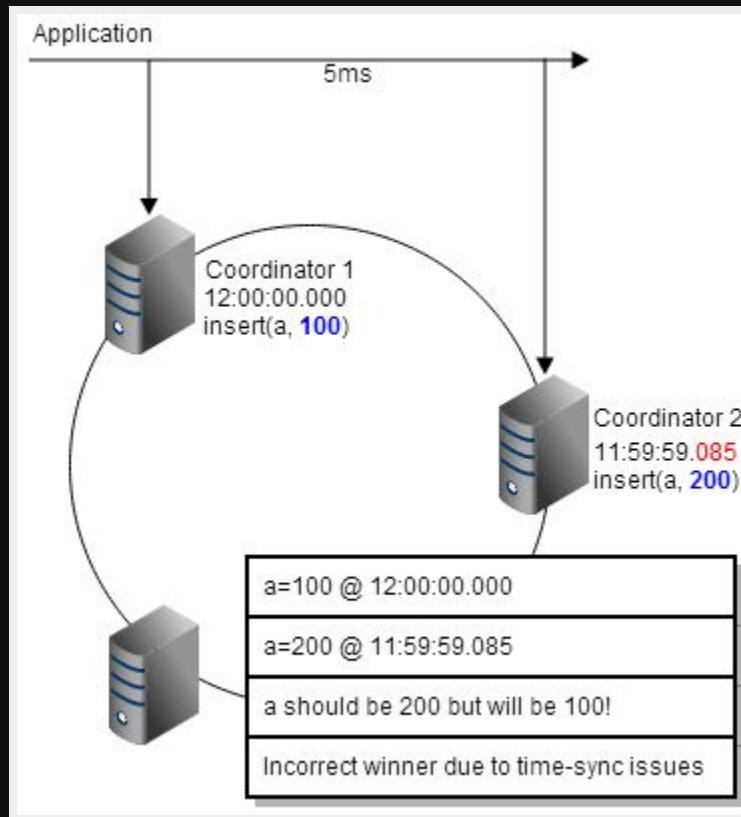
Does QUORUM read and
QUORUM write **guarantee**
consistency?

No - it gives strong consistency

NTP must be set-up correctly

Know your flow of data

Model as idem-potent-ly as possible



question 3

Will C* solve all your IT issues?

No - but it will help solve a lot of your problems

Performance and availability are **tunable** based on needs and costs

Consistency is solved by **design** and knowing the **flow** of data

Going from SQL to noSQL is a **paradigm** shift

Skills are not readily available

Some issues are better solved by **relational** databases

Thank you
...graceful bow

Contact
christopher.reedijk@ing.nl
gary.stewart@ing.nl

